INKLOOM AI

Intelligent Brand Governance Platform

ROUND 2 · THE BUILD

Working prototype for pre-publication brand content review — explainable checks, rule-grounded findings, and mandatory human approval.

Team name: [ENTER TEAM NAME]
Team leader: [FULL NAME] Email: [TEAM LEADER EMAIL]
Team members: [ENTER FULL NAMES]

Prepared for the We Code Coder hackathon · Round 2 submission

1. Executive Summary

Inkloom AI is a focused brand-governance assistant that reviews text-based marketing content (social posts, ads, short campaign copy) against an explicit library of brand rules and approved facts before publication. Every finding is labelled as rule-based or AI-assisted, linked to a guideline where possible, and never auto-publishes. A human reviewer retains final approval.

In Round 2 we delivered a complete end-to-end working prototype: brand rule library, draft submission, hybrid checker, explainable findings with suggestions, accept/reject workflow, human approval states, and revision history.

2. Completed Features

- Brand workspace with seeded sample guidelines (terminology, prohibited claims, tone, preferred language).
- Draft submission with channel selection (social / ad / campaign / email).
- Deterministic rule engine for exact prohibited phrases and claim patterns.
- AI-assisted heuristics for tone, message length, and brand-value language (clearly labelled).
- Severity ranking (critical → low) and confidence scores on every finding.
- Evidence-linked explanations that reference the matched guideline or pattern.
- Correction suggestions with one-click bulk apply.
- Per-finding Accept / Reject controls.
- Full human approval workflow: Draft → Under Review → Approved / Changes Requested.
- Revision history stored for every meaningful action.
- Dashboard with live counts and recent drafts.
- Add new guidelines from the UI (category, severity, exact-match flag).

3. Technology & Implementation

Backend: Python 3 + FastAPI. SQLAlchemy + SQLite for persistence. Modular checker that runs deterministic patterns first, then lighter heuristics labelled as AI-assisted. All findings returned as structured JSON.

Frontend: Lightweight single-page application (HTML/CSS/JS) for maximum reliability and zero build-step friction. Modern dark UI with clear severity badges, suggestion cards, and status pills.

Key technical decision: We deliberately keep the AI component transparent. Hard violations come from explicit rules; softer judgements are marked “AI-assisted” with lower confidence and always require human confirmation. This matches the Round 1 promise that the model is never treated as the brand standard.

Layer	Choice	Reason
Frontend	Vanilla JS + CSS	Reliable, fast, judge-friendly demo
Backend API	Python + FastAPI	Lightweight, typed, excellent docs
Validation	Rules + heuristics	Deterministic core + transparent AI layer
Storage	SQLite	Zero-config prototype persistence
Version control	Git	Shareable repository

4. Changes from the Round 1 Plan

- Frontend implemented as a polished static SPA instead of full React/Vite to guarantee a working demo under constrained runtime environments while preserving every planned screen and interaction.
- SQLite database path made environment-aware. Seed data provides an immediately usable demo workspace.
- AI layer deliberately kept hybrid and transparent (exact rules first) rather than calling an external LLM by default — this improves reliability, cost, and auditability for the prototype while still demonstrating the intended architecture.

All original workflow steps and value propositions remain intact.

5. Three Major Improvements for Round 3

1. Grounded LLM integration

Replace the current heuristics with a properly constrained LLM call that receives only the approved guideline set + the draft. Force structured JSON output. Keep the deterministic rule engine as the authoritative source for hard violations. This raises finding quality while preserving explainability and human control.

2. Visual brand checks

Allow upload of logo and colour-palette reference assets. Perform basic image analysis (dominant colours, logo presence) so that image-containing content can also be flagged for off-brand visuals — turning a Round 1 stretch goal into a core capability.

3. Recurring-issue analytics for brand managers

Aggregate findings across many drafts to surface the most frequent violation categories, unclear guidelines, and potential training needs. Provide a simple exportable report so the platform moves from one-off checks to continuous brand-governance insight.

6. How Judges Can Evaluate

1. Open the frontend and click “Load Sample Draft” then “Run Brand Checks”.
2. Inspect the findings — critical safety and guarantee claims should appear first, clearly labelled rule-based.
3. Accept or reject findings, optionally apply suggestions, then submit for review and approve.
4. Visit Brand Rules to see or add guidelines; they immediately affect subsequent checks.
5. Observe that nothing is ever auto-published and every status change is recorded.

Inkloom AI keeps brand governance proactive, explainable and collaborative. The Round 2 prototype delivers the complete text-review path promised in Round 1 and sets a clear foundation for the three substantial improvements planned for Round 3.

Before final submission: replace all bracketed team placeholders, verify the working demo link and repository link, and confirm the Google Form fields match this document.